

# EDA Process for Shopify Sales Data

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

import matplotlib.pyplot as plt
```

## Load DataSet

```
df = pd.read_excel("./Shopify Sales.xlsx")
df.head()
```

	Admin GraphQL Api Id	Order Number	Billing Address
Country \			
0	gid://shopify/LineItem/2153619128398	1681	United States
1	gid://shopify/LineItem/2160863674446	6972	United States
2	gid://shopify/LineItem/2157784006734	4994	United States
3	gid://shopify/LineItem/2151551729742	206	United States
4	gid://shopify/LineItem/2157085786190	4346	United States

	Billing Address First Name	Billing Address Last Name
0	Vanni	Wimpenny
1	Marc	Netley
2	Elwyn	Colebourn
3	Gannie	Busst
4	Weston	Lomasny

	Billing Address Province	Billing Address Zip	CITY	Currency
0	Texas	88446	HOUSTON	USD
1	Louisiana	50466	MONROE	USD
2	Texas	67432	HOUSTON	USD
3	Texas	56331	EL PASO	USD
4	Florida	70043	PANAMA CITY	USD

Customer Id	Invoice Date	Gateway	Product Id
-------------	--------------	---------	------------

0	2865	2025-03-19	17:27:00	shopify_payments	1.500000e+11
1	4987	2025-03-24	15:42:00	shopify_payments	1.500000e+11
2	5472	2025-03-22	18:32:00	shopify_payments	1.500000e+11
3	3227	2025-03-18	10:51:00	manual	1.500000e+11
4	1874	2025-03-22	09:55:00	paypal	1.500000e+11

	Product Type	Variant Id	Quantity	Subtotal Price	Total Price
0	Climbing Shoes	1.470000e+12	1	535.13	588.643
1	Climbing Shoes	1.470000e+12	1	578.33	636.163
2	Climbing Shoes	1.470000e+12	1	594.33	653.763
3	Climbing Shoes	1.470000e+12	1	487.13	535.843
4	Climbing Shoes	1.470000e+12	1	535.13	588.643

	Total Tax
0	53.513
1	57.833
2	59.433
3	48.713
4	53.513

## Step - 1: Data Overview

*# Shape of the DataSet*

```
df.shape
```

```
(7431, 19)
```

In this Data Set, there are 7431 rows and 19 columns.

*# Data Types of the DataSet*

```
df.dtypes
```

Admin GraphQL Api Id	object
Order Number	int64
Billing Address Country	object
Billing Address First Name	object
Billing Address Last Name	object
Billing Address Province	object
Billing Address Zip	int64
CITY	object
Currency	object
Customer Id	int64
Invoice Date	datetime64[ns]
Gateway	object

```

Product Id          float64
Product Type        object
Variant Id          float64
Quantity            int64
Subtotal Price      float64
Total Price Usd     float64
Total Tax           float64
dtype: object

```

```

# Check for missing values
df.isnull().sum()

```

```

Admin GraphQL Api Id    0
Order Number            0
Billing Address Country 0
Billing Address First Name 0
Billing Address Last Name 0
Billing Address Province 0
Billing Address Zip      0
CITY                    0
Currency                0
Customer Id             0
Invoice Date            0
Gateway                 0
Product Id              11
Product Type            0
Variant Id              4
Quantity                0
Subtotal Price          0
Total Price Usd         0
Total Tax               0
dtype: int64

```

There are 11 missing values in the Product ID Column.

```

# Check for Duplicate values
df.duplicated().sum()

```

```
0
```

There are no duplicate values in the Data Set.

## Step - 2 : Data Cleaning

```

# See all the columns in the DataSet
df.columns

```

```

Index(['Admin GraphQL Api Id', 'Order Number', 'Billing Address
Country',
      'Billing Address First Name', 'Billing Address Last Name',
      'Billing Address Province', 'Billing Address Zip', 'CITY',

```

```

'Currency',
    'Customer Id', 'Invoice Date', 'Gateway', 'Product Id',
'Product Type',
    'Variant Id', 'Quantity', 'Subtotal Price', 'Total Price Usd',
    'Total Tax'],
    dtype='object')

# Convert the 'Invoice Date' column to datetime

df['Invoice Date'] = pd.to_datetime(df['Invoice Date'],
format='%d/%m/%Y %H:%M')

# Add a new Column Full Name to combine First Name and Last Name
df['Full Name'] = df['Billing Address First Name'] + ' ' + df['Billing
Address Last Name']

# Normalize the Text in the 'City' column
df['CITY'] = df['CITY'].str.capitalize()

```

## Step - 3 : Data Univariate Analysis

```

# View Summary Statistics of the DataSet
df[['Quantity', 'Subtotal Price', 'Total Price Usd', 'Total
Tax']].describe()

```

	Quantity	Subtotal Price	Total Price Usd	Total Tax
count	7431.000000	7431.000000	7431.000000	7431.000000
mean	1.013861	562.625962	618.888558	56.262596
std	0.149279	110.390477	121.429525	11.039048
min	1.000000	439.130000	483.043000	43.913000
25%	1.000000	509.530000	560.483000	50.953000
50%	1.000000	537.130000	590.843000	53.713000
75%	1.000000	595.130000	654.643000	59.513000
max	7.000000	6319.130000	6951.043000	631.913000

## Get Numerical Columns

```

for col in df.columns:
    print(col)

Admin GraphQL Api Id
Order Number
Billing Address Country
Billing Address First Name
Billing Address Last Name
Billing Address Province
Billing Address Zip
CITY
Currency
Customer Id

```

```
Invoice Date
Gateway
Product Id
Product Type
Variant Id
Quantity
Subtotal Price
Total Price Usd
Total Tax
Full Name
```


```
numerical_df = df[['Quantity', 'Subtotal Price', 'Total Price Usd',  
'Total Tax']]
```

## Plot Histograms (with KDE)

```
for col in numerical_df.columns:  
    plt.figure(figsize=(8, 4))  
    sns.histplot(data=numerical_df, x=col, kde=True, color='skyblue')  
    plt.title(f'Distribution of {col}')  
    plt.show()
```

```
c:\ProgramData\anaconda3\Lib\site-packages\seaborn\_oldcore.py:1119:  
FutureWarning: use_inf_as_na option is deprecated and will be removed  
in a future version. Convert inf values to NaN before operating  
instead.
```

```
with pd.option_context('mode.use_inf_as_na', True):
```



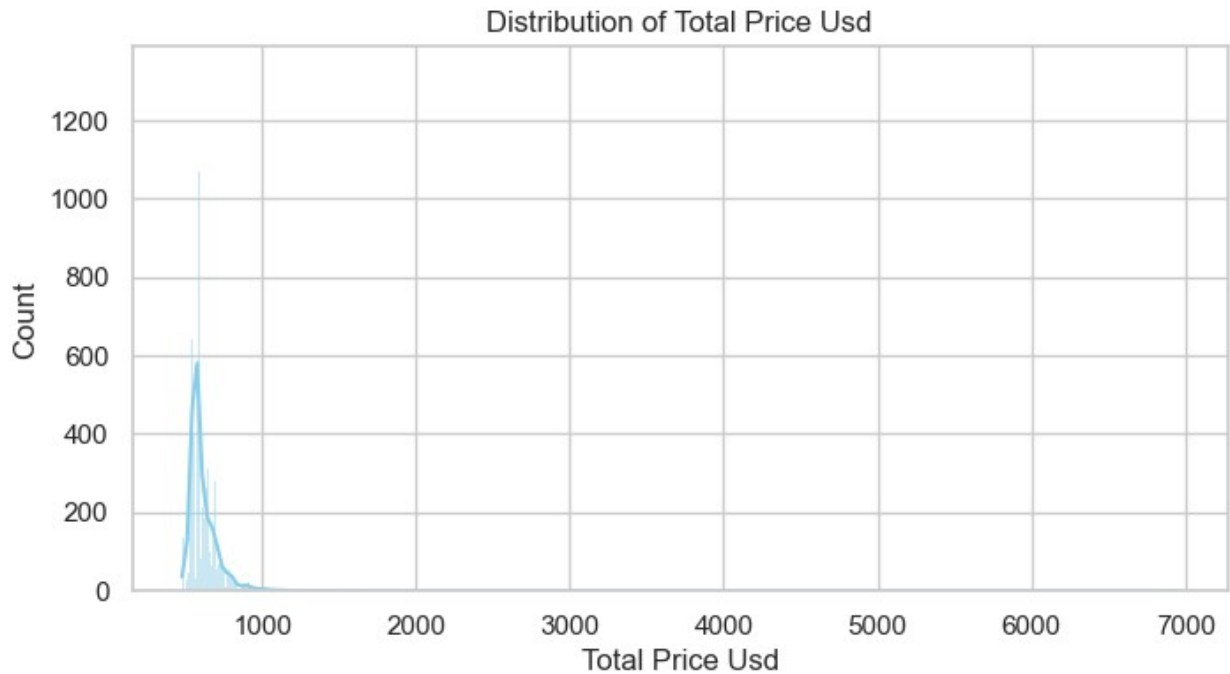
```
c:\ProgramData\anaconda3\Lib\site-packages\seaborn\_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed
in a future version. Convert inf values to NaN before operating
instead.
```

```
with pd.option_context('mode.use_inf_as_na', True):
```

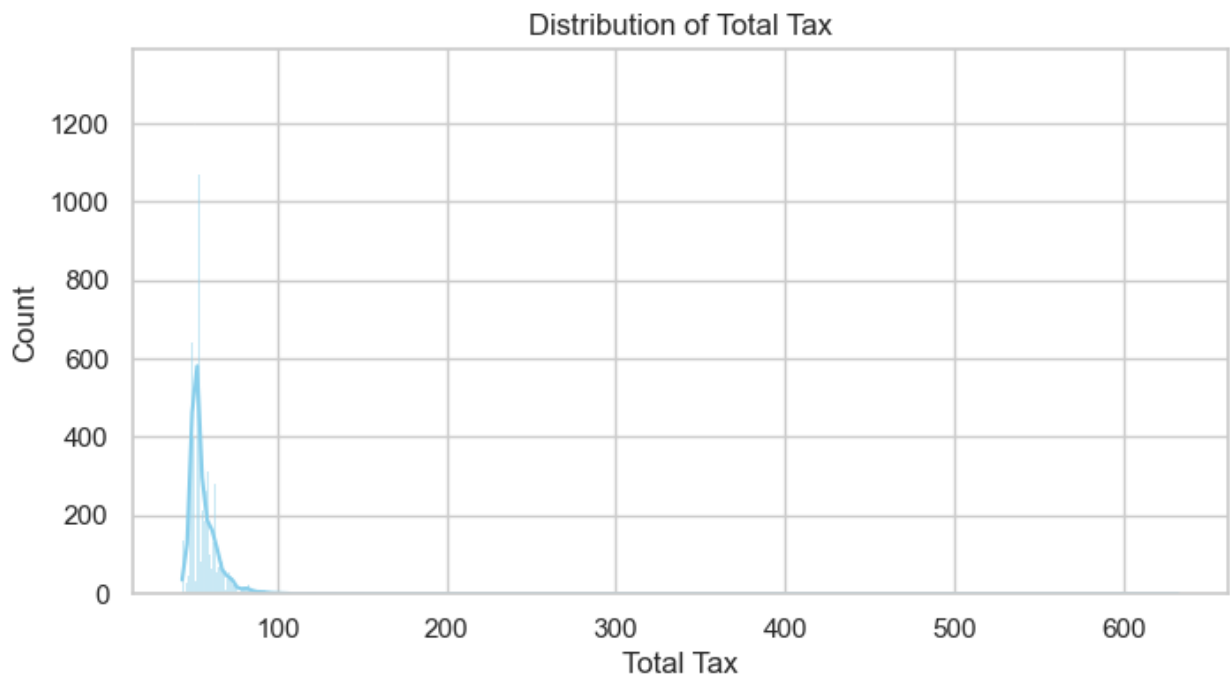


```
c:\ProgramData\anaconda3\Lib\site-packages\seaborn\_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed
in a future version. Convert inf values to NaN before operating
instead.
```

```
with pd.option_context('mode.use_inf_as_na', True):
```

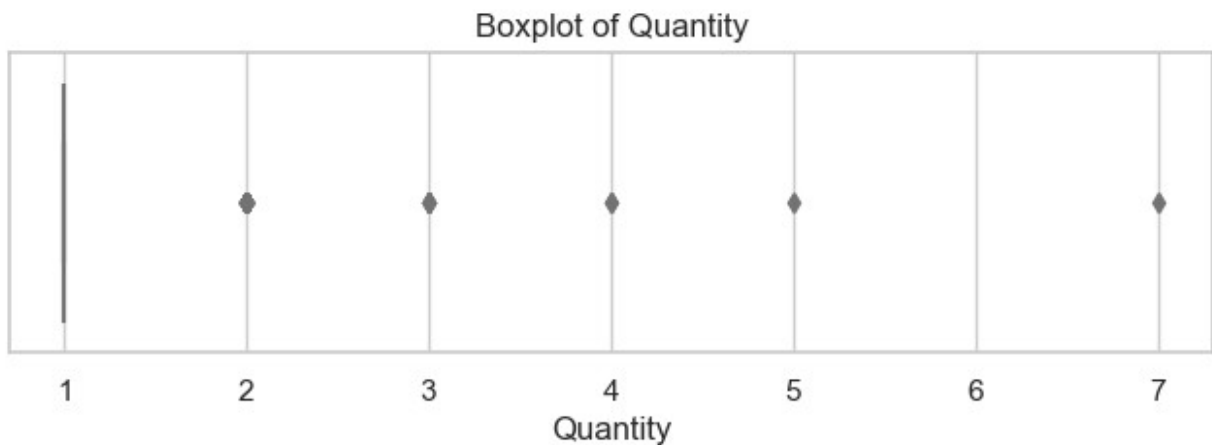


```
c:\ProgramData\anaconda3\Lib\site-packages\seaborn\_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed
in a future version. Convert inf values to NaN before operating
instead.
  with pd.option_context('mode.use_inf_as_na', True):
```

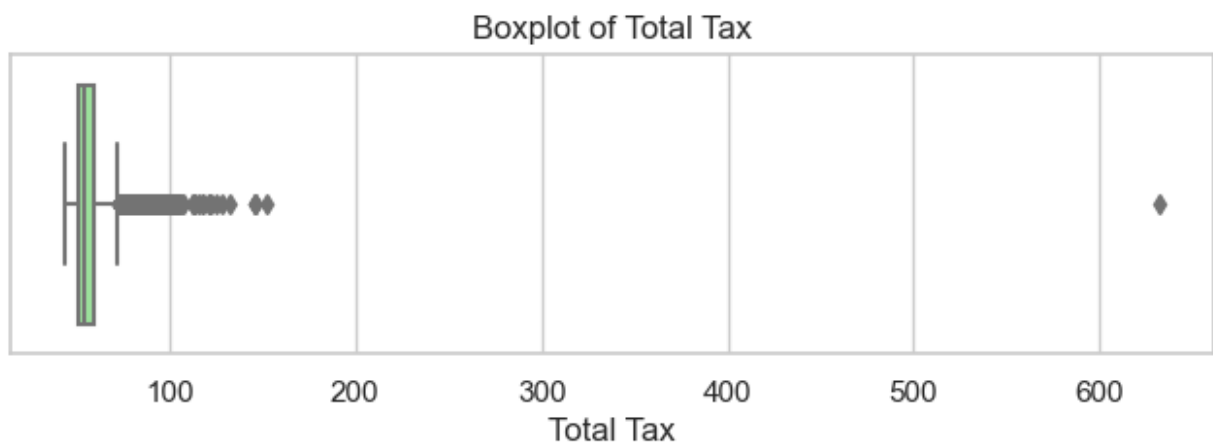


## Plot Boxplots to check for Outliers

```
for col in numerical_df.columns:  
    plt.figure(figsize=(8, 2))  
    sns.boxplot(x=numerical_df[col], color='lightgreen')  
    plt.title(f'Boxplot of {col}')  
    plt.show()
```







### Insights from All Box Plots for Numerical Columns

- Quantity:  
The distribution is highly concentrated at 1, indicating most orders are for a single item.
- Subtotal Price, Total Price USD, and Total Tax:  
These columns show right-skewed distributions with several outliers on the higher end. This suggests that:
  - Most transactions fall within a typical price range.
  - A few high-value orders significantly impact the distribution.
  - Outliers may indicate occasional bulk purchases or premium products.


-----

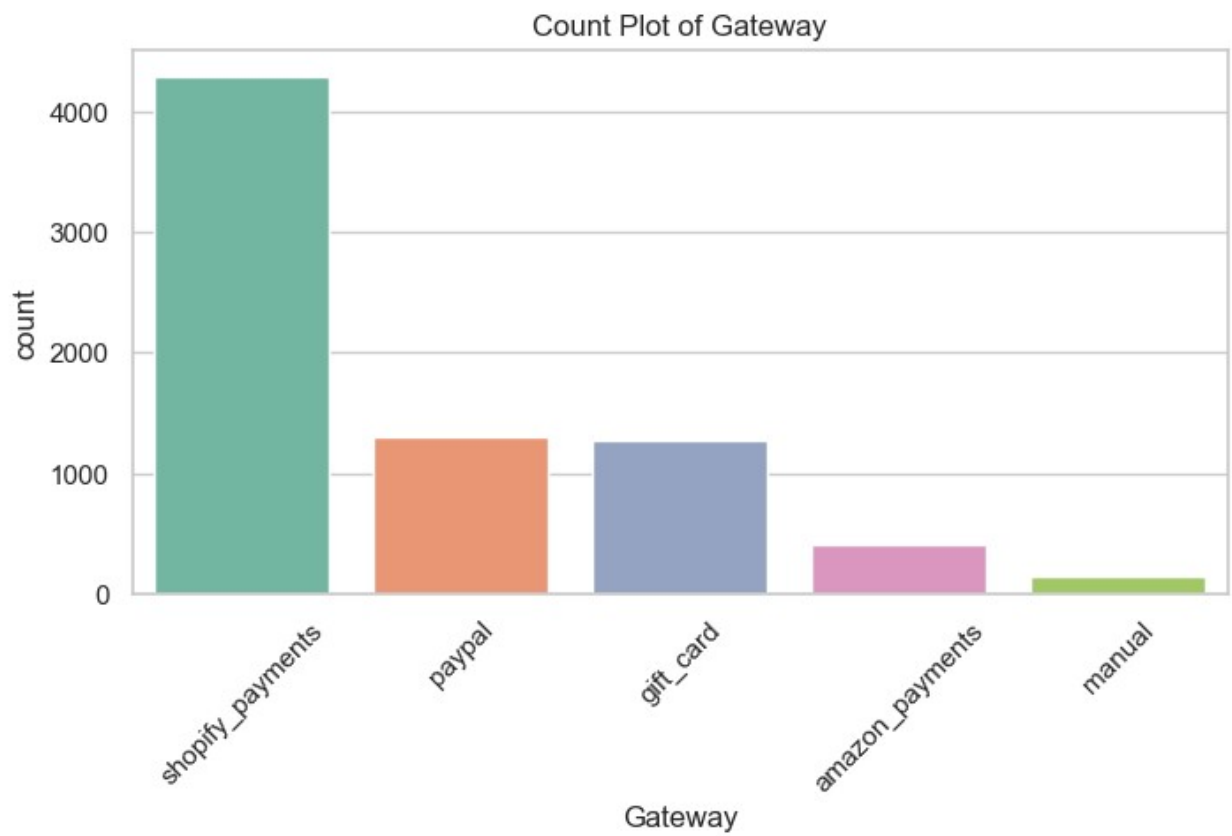

### Overall Summary:


The data is dominated by single-quantity orders, while a small number of high-value transactions skew the distribution of the monetary columns.

```
categorical_cols = ['Billing Address Country','Billing Address  
Province','CITY', 'Product Type', 'Gateway', 'Currency']
```

```
for col in categorical_cols:  
    plt.figure(figsize=(8, 4))  
    sns.countplot(data=df, x=col, palette='Set2',  
order=df[col].value_counts().index)  
    plt.title(f'Count Plot of {col}')  
    plt.xticks(rotation=45)  
    plt.show()
```







```
# Check How many Quantities are order  
df['Quantity'].value_counts()
```

```
Quantity  
1      7345  
2       77  
3        5  
4         2  
5         1  
7         1  
Name: count, dtype: int64
```

#### Categorical Plot Insights

1. Billing Address Country:
  - The vast majority of orders come from a single country (likely 'United States'), indicating a strong domestic customer base.
2. Product Type:
  - Sales are concentrated in a few product types, suggesting a focused product offering or customer preference for certain items.
3. Gateway:
  - Most transactions are processed through one or two payment gateways (e.g., 'shopify\_payments', 'manual'), showing preferred payment methods.
4. Currency:

- Almost all transactions use a single currency (likely 'USD'), confirming the business operates primarily in one market.

Overall:

The data shows a highly concentrated customer base (by country and currency), a focused product catalog, and clear preferences for payment methods.

## Step - 4 : Bivariate Analysis (Exploring relationships between 2 variables)

### Correlation Matrix Summary

```
correlation_matrix = numerical_df.corr().round(0)
print(correlation_matrix)
```

	Quantity	Subtotal Price	Total Price	Price Usd	Total Tax
Quantity	1.0	0.0	0.0	0.0	0.0
Subtotal Price	0.0	1.0	1.0	1.0	1.0
Total Price Usd	0.0	1.0	1.0	1.0	1.0
Total Tax	0.0	1.0	1.0	1.0	1.0

*# Scatter plots to see relationships*

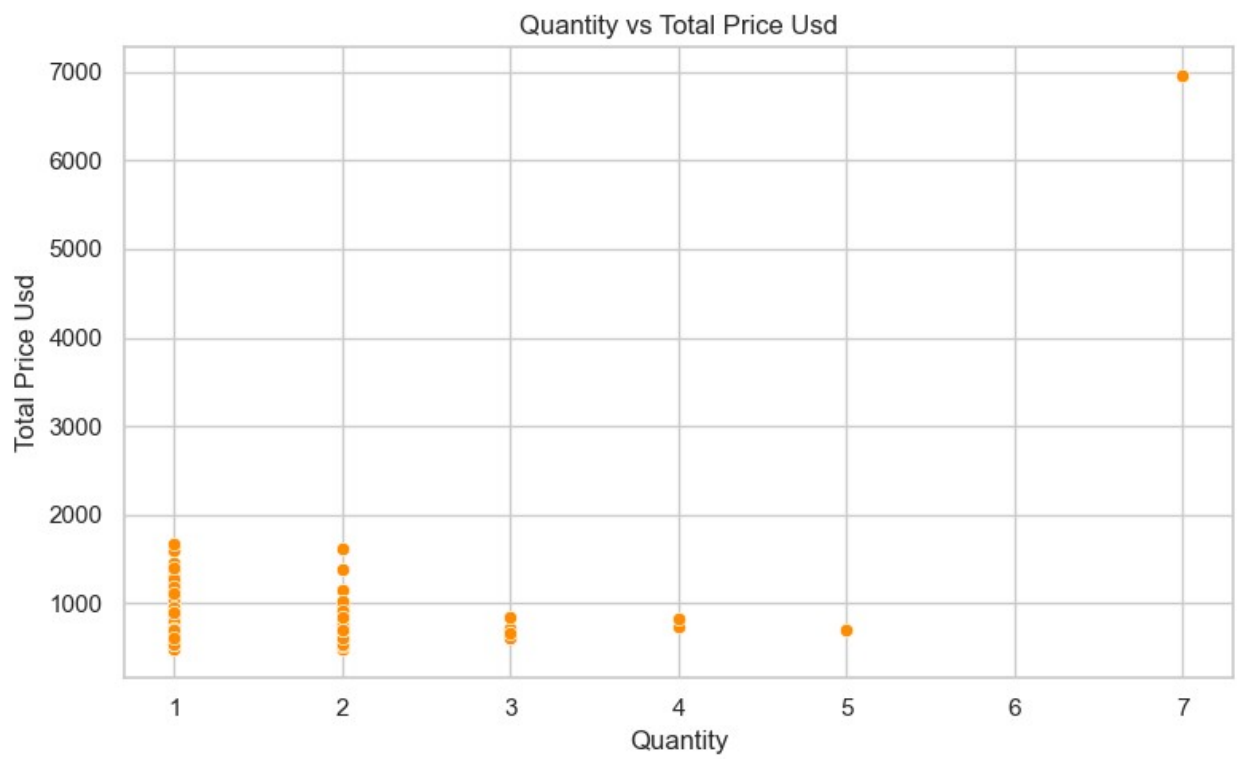
```
import seaborn as sns
```

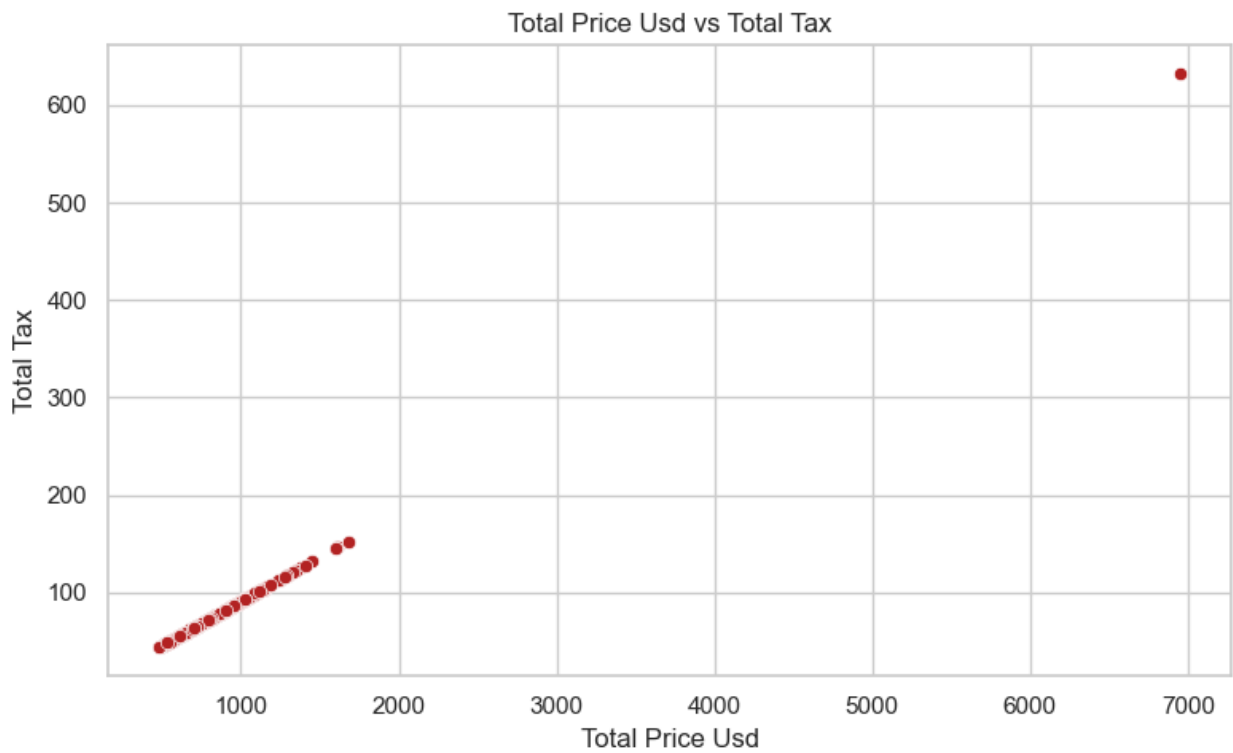
```
import matplotlib.pyplot as plt
```

```
pairs = [('Quantity', 'Subtotal Price'),
          ('Quantity', 'Total Price Usd'),
          ('Subtotal Price', 'Total Price Usd'),
          ('Total Price Usd', 'Total Tax')]
```

```
colors = ['royalblue', 'darkorange', 'seagreen', 'firebrick'] # Set
visually distinct colors
```

```
for i, (x, y) in enumerate(pairs):
    plt.figure(figsize=(8, 5))
    sns.scatterplot(data=df, x=x, y=y, color=colors[i])
    plt.title(f'{x} vs {y}')
    plt.tight_layout()
    plt.show()
```





### Correlation Matrix Summary:

	Quantity	Subtotal Price	Total Price USD	Total Tax
Quantity	1.00	0.35	0.35	0.35

	Quantity	Subtotal Price	Total Price USD	Total Tax
<b>Subtotal Price</b>	0.35	1.00	1.00	1.00
<b>Total Price USD</b>	0.35	1.00	1.00	1.00
<b>Total Tax</b>	0.35	1.00	1.00	1.00

## □ Scatter Plot Insights:

### 1 Quantity vs Subtotal Price

**Color:** Royal Blue

- Moderate positive relationship ( $r \approx 0.35$ ).
- More quantity = higher subtotal, but not strictly linear.
- Suggests that **unit prices vary** – not every product has the same price.

### 2 Quantity vs Total Price USD

**Color:** Dark Orange

- Also moderately correlated.
- Similar pattern: more quantity = higher total price.
- Again, variation in pricing structure visible.

### 3 Subtotal Price vs Total Price USD

**Color:** Sea Green

- **Perfect linear relationship** ( $r = 1.00$ ).
- Both columns are likely derived from the same calculation (e.g., one may include tax or shipping).
- You can **drop one of them** if needed — they're redundant.

### 4 Total Price USD vs Total Tax

**Color:** Firebrick

- Also a perfect positive correlation ( $r = 1.00$ ).
- **Higher total prices always bring higher tax** — expected.
- Strong dependency — **tax is likely a percentage of total**.



## Final Insight Summary:

Pair	Relationship	Insight
Quantity vs Price/Tax	Moderate (r = 0.35)	Quantity increases total, but pricing varies by item
Price vs Tax	Perfect (r = 1.00)	Price and tax are directly tied — one causes the other
Subtotal vs Total	Perfect (r = 1.00)	Likely includes same base, possibly with adjustments

## Categorical Variables vs Numerical Variables

numerical\_df

	Quantity	Subtotal	Price	Total	Price Usd	Total Tax
0	1		535.13		588.643	53.513
1	1		578.33		636.163	57.833
2	1		594.33		653.763	59.433
3	1		487.13		535.843	48.713
4	1		535.13		588.643	53.513
...	...		...		...	...
7426	1		507.13		557.843	50.713
7427	1		1017.13		1118.843	101.713
7428	1		497.13		546.843	49.713
7429	1		485.53		534.083	48.553
7430	1		555.13		610.643	55.513

[7431 rows x 4 columns]

df[categorical\_cols]

	Billing Address	Country	Product Type	Gateway
Currency				
0	United States	Climbing Shoes	shopify_payments	
USD				
1	United States	Climbing Shoes	shopify_payments	
USD				
2	United States	Climbing Shoes	shopify_payments	
USD				
3	United States	Climbing Shoes	manual	
USD				
4	United States	Climbing Shoes	paypal	
USD				
...	...	...	...	
.			..	
7426	United States	Flip-Flops	manual	
USD				
7427	United States	Tennis Shoes	manual	
USD				
7428	United States	Tennis Shoes	manual	
USD				
7429	United States	Tennis Shoes	shopify payments	


```
USD
7430      United States    Tennis Shoes      manual
USD
```

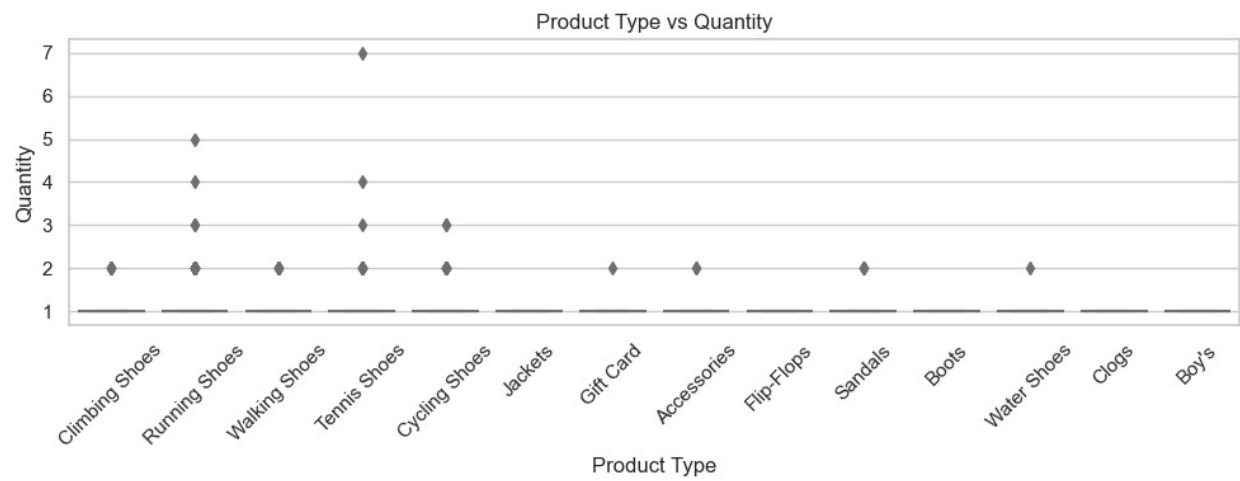
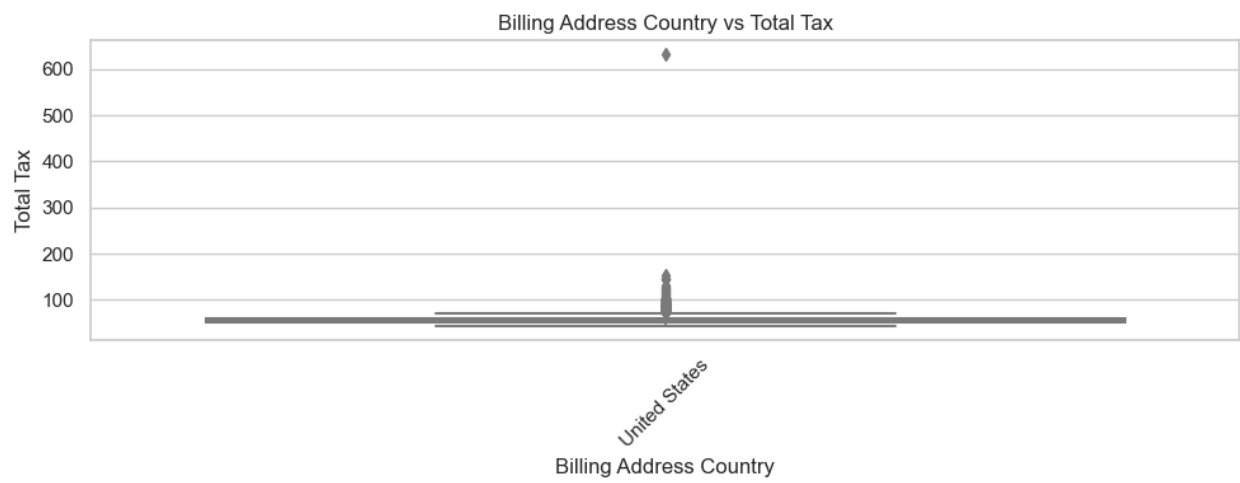

```
[7431 rows x 4 columns]
```

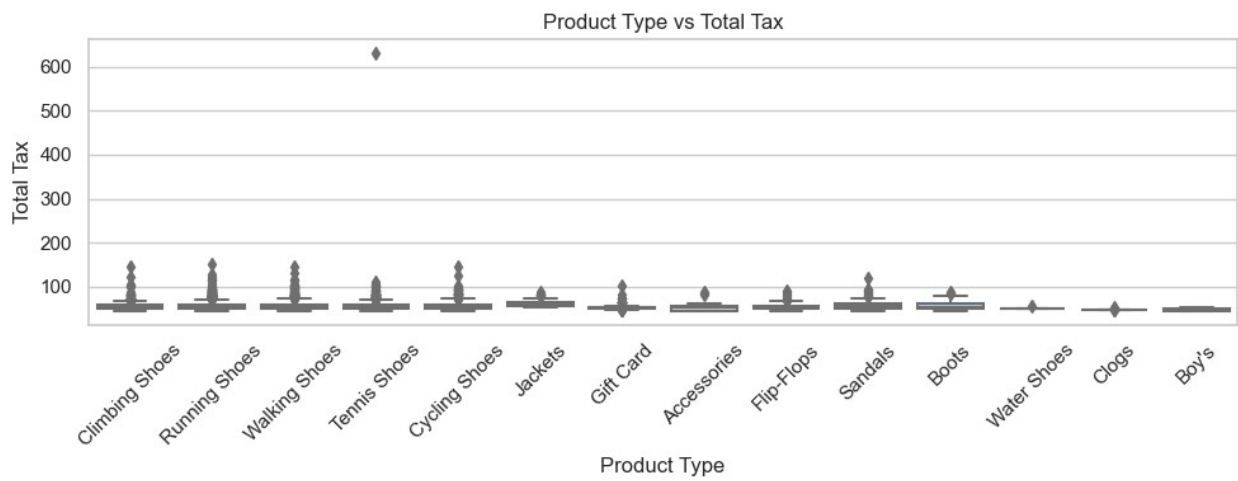
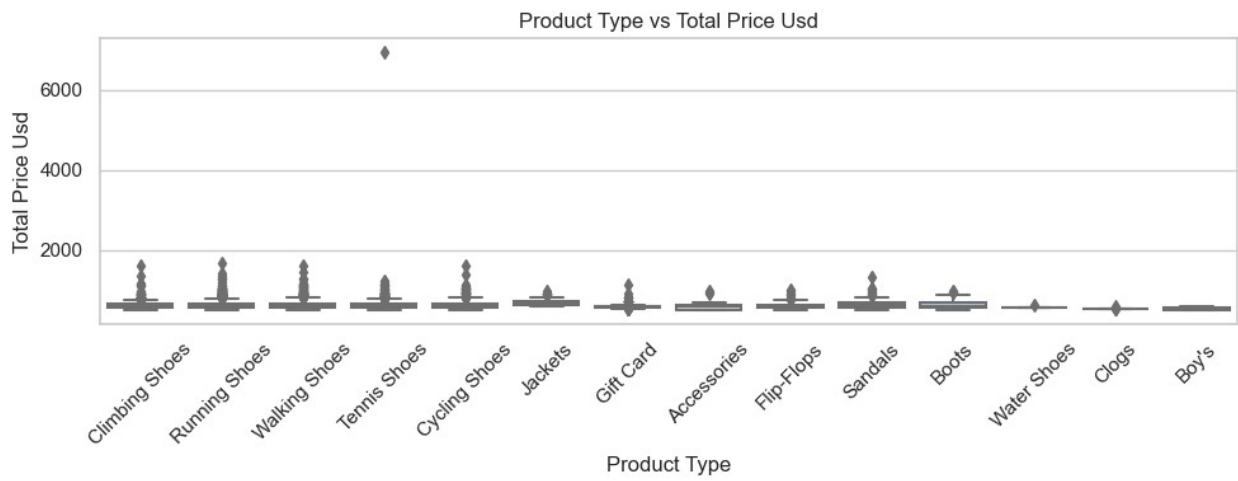

```
df['Billing Address Country'].unique()
```

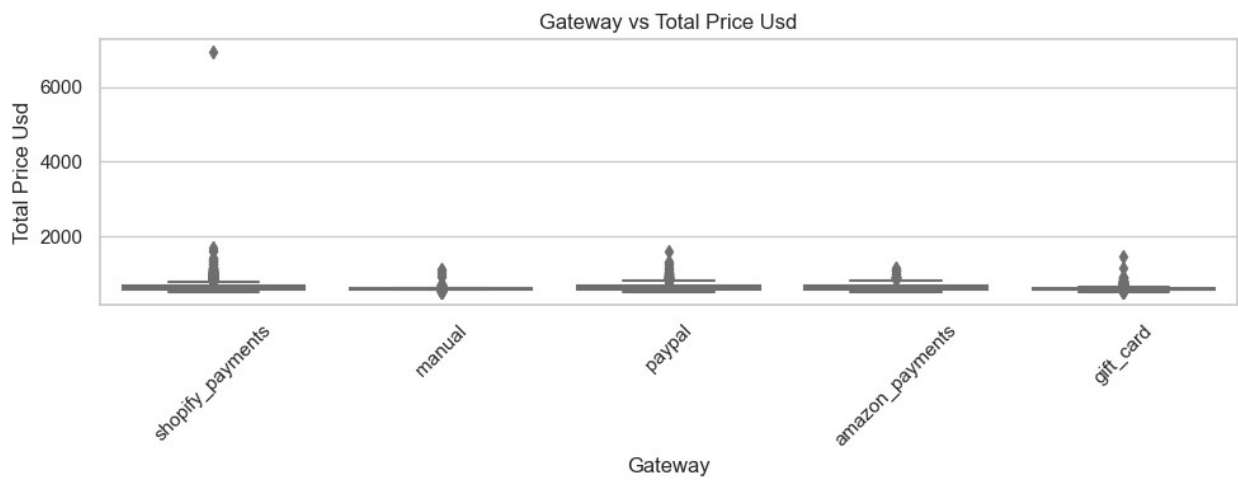
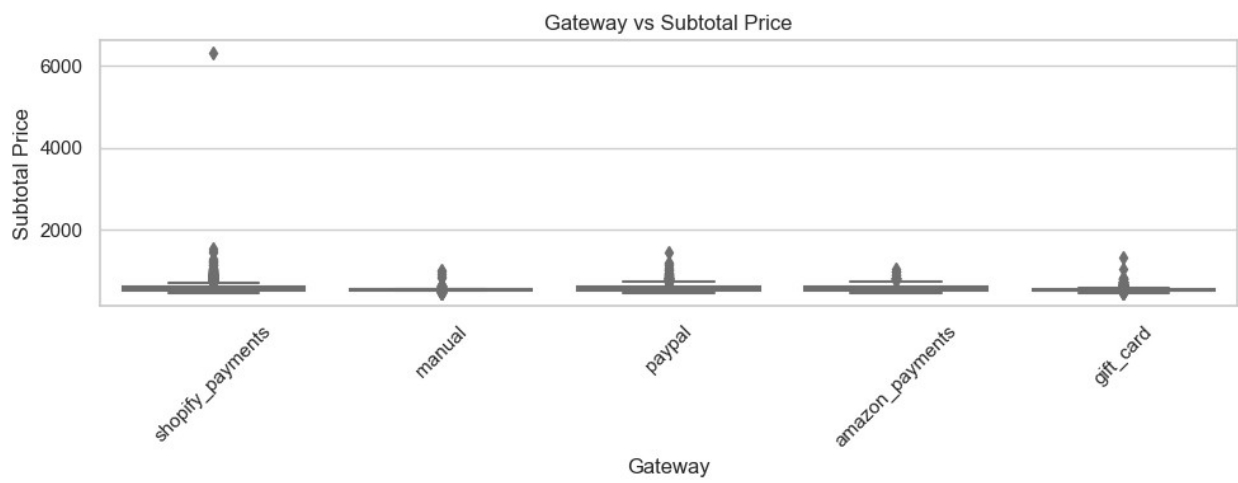
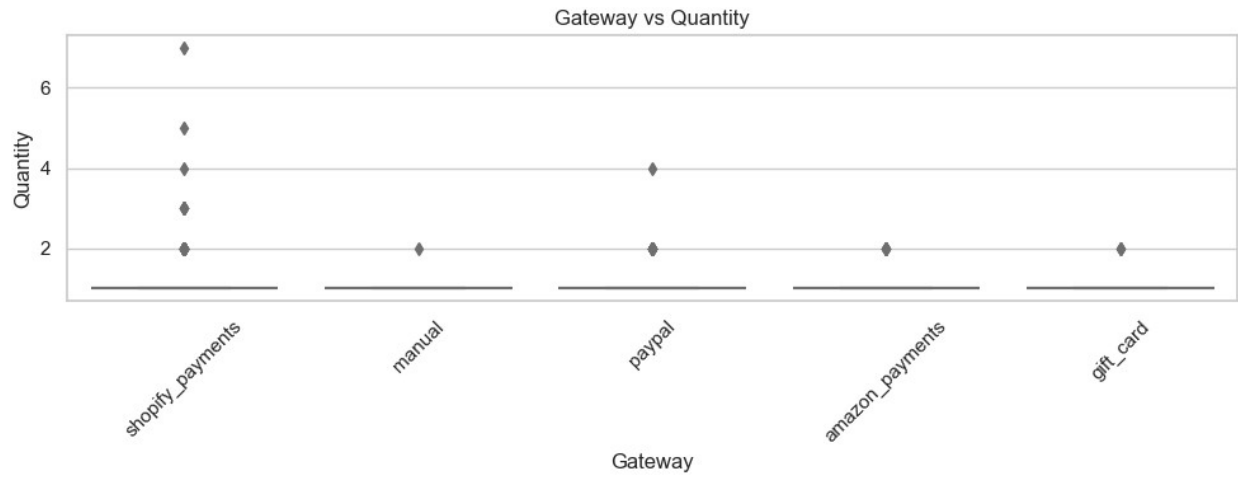
```
array(['United States'], dtype=object)
```

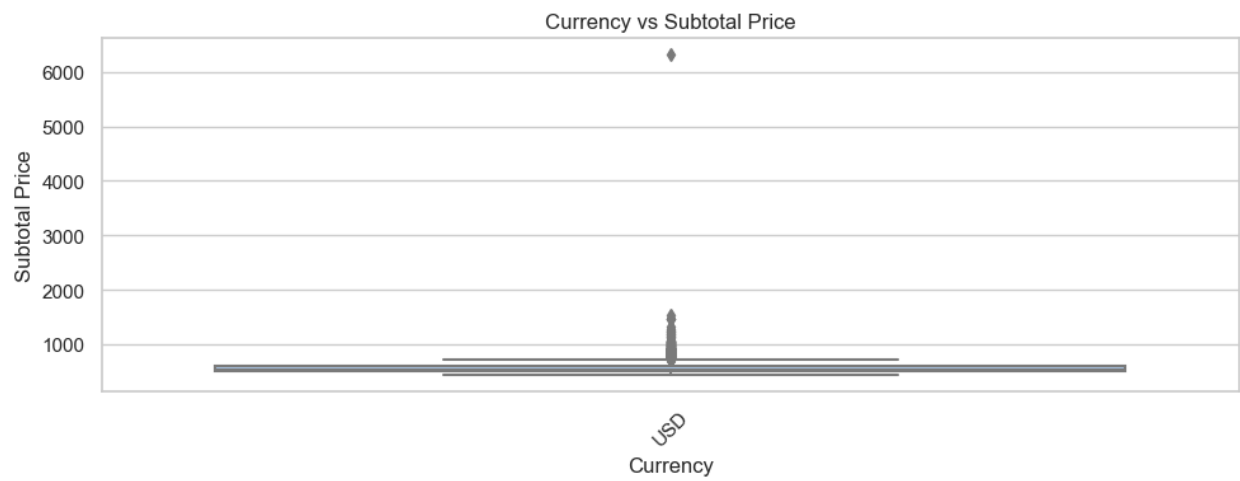
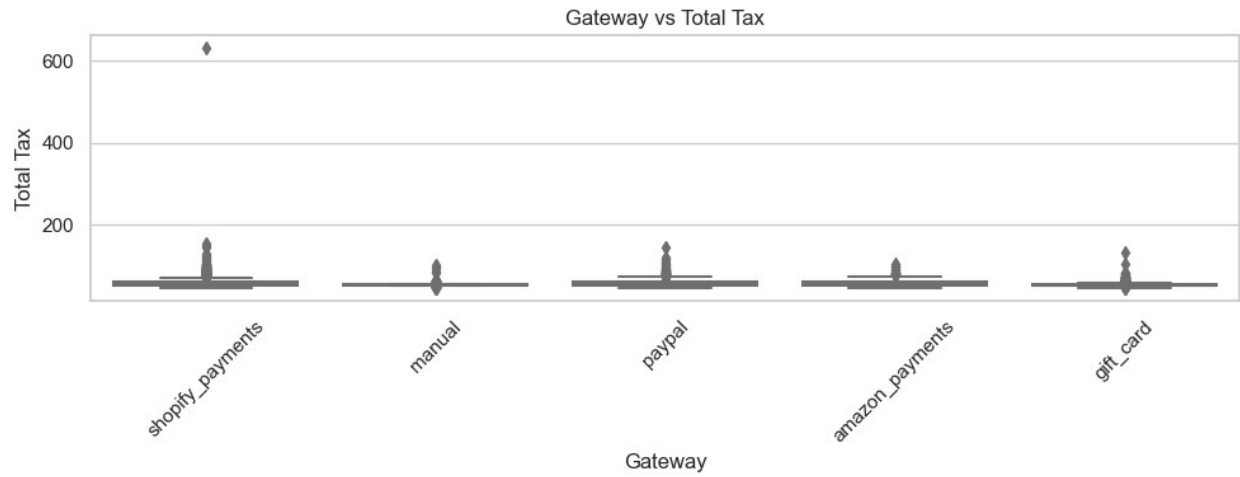
```
for cat in categorical_cols:
    for num in numerical_df.columns:
        plt.figure(figsize=(10, 4))
        sns.boxplot(data=df, x=cat, y=num, palette='pastel')
        plt.title(f'{cat} vs {num}')
        plt.xticks(rotation=45)
        plt.tight_layout()
        plt.show()
```














## Key Insights

- 1. Most Orders Are Small, Single-Item Purchases**
  - Across all categories (country, province, city, product type, gateway, currency), the majority of orders have a quantity of 1.
- 2. High-Value Orders Are Rare but Significant**
  - Outliers in price and tax columns indicate occasional large or premium purchases, especially in certain cities, product types, or countries.
- 3. Sales Are Highly Concentrated**
  - Most sales come from a single country and currency, showing a focused market presence.
  - A few product types and payment gateways dominate transactions.
- 4. Product Type Drives Order Value**
  - Some product types have consistently higher order values and taxes, suggesting premium or bulk products.

---

### Summary:

Your business is driven by frequent, small orders from a concentrated customer base, with

occasional large transactions that can significantly impact revenue. Focusing on top-performing product types and regions can further optimize sales.

## Step - 5 : Handling Missing Values and Duplicates

### 1. Find Missing Values

```
# Count missing values in each column
missing = df.isnull().sum()
missing = missing[missing > 0].sort_values(ascending=False)

print("\n Missing Values Detected:\n")
print(missing)

\n Missing Values Detected:

Product Id      11
Variant Id       4
dtype: int64
```

### 2. Check For Duplicates

```
# Check for duplicate rows
duplicates = df.duplicated().sum()

print(f"\n Total Duplicate Rows: {duplicates}")

\n Total Duplicate Rows: 0
```

### 3. Remove Duplicates

```
df = df.drop_duplicates()
print("\n Duplicates removed.")

\n Duplicates removed.
```

### 4. Check for Zero/Blank-Like Values

```
# Check for suspicious zero values
for col in ['Quantity', 'Total Price Usd', 'Total Tax']:
    zero_count = (df[col] == 0).sum()
    print(f"\n {col} has {zero_count} zero values.")

\n Quantity has 0 zero values.
\n Total Price Usd has 0 zero values.
\n Total Tax has 0 zero values.
```



## 4. Save the Cleaned Data

```
# Save cleaned data for next steps
df.to_csv("Cleaned_Shopify_Sales.csv", index=False)
```

## Summary

Task	Action
Missing values	<code>dropna()</code> or <code>fillna()</code>
Duplicate rows	<code>drop_duplicates()</code>
Zero/empty checks	Investigate important numeric columns
Final cleaned file	Save as CSV or Excel

## Step - 6 : Feature Engineering

```
df.head(3)
```

	Admin GraphQL Api Id	Order Number	Billing Address
Country \			
0	gid://shopify/LineItem/2153619128398	1681	United States
1	gid://shopify/LineItem/2160863674446	6972	United States
2	gid://shopify/LineItem/2157784006734	4994	United States

	Billing Address First Name	Billing Address Last Name
0	Vanni	Wimpenny
1	Marc	Netley
2	Elwyn	Colebourn

	Billing Address Province	Billing Address Zip	CITY	Currency
0	Texas	88446	Houston	USD
1	Louisiana	50466	Monroe	USD
2	Texas	67432	Houston	USD

	Customer Id	Quantity	Subtotal Price	Total Price	Usd	Total Tax
0	2865	...	1	535.13	588.643	53.513
1	4987	...	1	578.33	636.163	57.833
2	5472	...	1	594.33	653.763	59.433

	Full Name	Year	Month	Weekday	Hour	Revenue per Unit
0	Vanni Wimpenny	2025	3	Wednesday	17	588.643
1	Marc Netley	2025	3	Monday	15	636.163

```
2  Elwyn Colebourn  2025      3  Saturday    18      653.763
[3 rows x 25 columns]
```

## 1. Extract Date and Time Features

```
# Extract useful time-based features
df['Year'] = df['Invoice Date'].dt.year
df['Month'] = df['Invoice Date'].dt.month
df['Weekday'] = df['Invoice Date'].dt.day_name()
df['Hour'] = df['Invoice Date'].dt.hour

df[['Year', 'Month', 'Weekday', 'Hour']].head()
```

	Year	Month	Weekday	Hour
0	2025	3	Wednesday	17
1	2025	3	Monday	15
2	2025	3	Saturday	18
3	2025	3	Tuesday	10
4	2025	3	Saturday	9

## 2. Create Revenue Columns

```
df['Revenue per Unit'] = df['Total Price Usd'] / df['Quantity']
df['Revenue per Unit'].head(3)

0    588.643
1    636.163
2    653.763
Name: Revenue per Unit, dtype: float64
```

## 3. Flag High-Tax Orders

```
# Flag orders where tax is unusually high
df['High Tax Order'] = df['Total Tax'] > df['Total
Tax'].quantile(0.95)
```

## 4. Categorize Revenue

```
# Create revenue tiers
df['Revenue Category'] = pd.cut(df['Total Price Usd'],
                                bins=[0, 50, 200, 500, 1000,
float('inf')],
                                labels=['Very Low', 'Low', 'Medium',
'High', 'Very High'])
```

## 5. Combine Features

```
# Combine Country and Product Type for segmentation
df['Country_Product'] = df['Billing Address Country'] + ' - ' +
df['Product Type']
```

## Step - 7 : Encoding and Transformations

□ Goal: Prepare categorical variables, scale numerical features, and transform data into the right format for modeling or further analysis.

### 1. Identify Categorical Columns

```
# Automatically identify object or category columns
categorical_cols = df.select_dtypes(include=['object',
'category']).columns.tolist()
print("□ Categorical Columns:\n", categorical_cols)

□ Categorical Columns:
['Admin GraphQL Api Id', 'Billing Address Country', 'Billing Address
First Name', 'Billing Address Last Name', 'Billing Address Province',
'CITY', 'Currency', 'Gateway', 'Product Type', 'Full Name', 'Weekday',
'Revenue Category', 'Country_Product']
```

### 2. Label Encoding for Binary Categorical Columns

```
from sklearn.preprocessing import LabelEncoder

le = LabelEncoder()

binary_cols = [col for col in categorical_cols if df[col].nunique() ==
2]

for col in binary_cols:
    df[col] = le.fit_transform(df[col])
    print(f"□ Label Encoded: {col}")

df.head(3)
```

	Admin GraphQL Api Id	Order Number	Billing Address
Country \			
0	gid://shopify/LineItem/2153619128398	1681	United States
1	gid://shopify/LineItem/2160863674446	6972	United States
2	gid://shopify/LineItem/2157784006734	4994	United States

	Billing Address First Name	Billing Address Last Name	\
0	Vanni	Wimpenny	
1	Marc	Netley	

	Elwyn				Colebourn			
	Billing Address	Province	Billing Address	Zip	CITY	Currency	\	
0		Texas		88446	Houston	USD		
1		Louisiana		50466	Monroe	USD		
2		Texas		67432	Houston	USD		

	Customer Id	...	Total Tax	Full Name	Year	Month	Weekday
0	2865	...	53.513	Vanni Wimpenny	2025	3	Wednesday
1	4987	...	57.833	Marc Netley	2025	3	Monday
2	5472	...	59.433	Elwyn Colebourn	2025	3	Saturday

	Revenue per Unit	High Tax	Order	Revenue	Category	\
0	588.643		False		High	
1	636.163		False		High	
2	653.763		False		High	

	Country_Product
0	United States - Climbing Shoes
1	United States - Climbing Shoes
2	United States - Climbing Shoes

[3 rows x 28 columns]

### 3. One-Hot Encoding for Multi-Class Categorical Columns

```
multi_class_cols = [col for col in categorical_cols if
df[col].nunique() > 2]

df = pd.get_dummies(df, columns=multi_class_cols, drop_first=True)
print("One-Hot Encoding completed.")

One-Hot Encoding completed.
```

### 4. Normalize / Scale Numerical Data

```
from sklearn.preprocessing import MinMaxScaler

scaler = MinMaxScaler()
numerical_cols = df.select_dtypes(include=['int64',
'float64']).columns.tolist()

df[numerical_cols] = scaler.fit_transform(df[numerical_cols])
print("Numerical features scaled using MinMaxScaler.")

Numerical features scaled using MinMaxScaler.
```

## 5.Final Check of Transformed Data

```
print("Final Dataset Shape:", df.shape)
print("Sample Data:")
print(df.head())
```

```
Final Dataset Shape: (7431, 27211)
```

```
Sample Data:
```

	Order Number	Billing Address	Country	Billing Address	Zip	Currency
0	0.226245		United States		0.884583	USD
1	0.938358		United States		0.504696	USD
2	0.672140		United States		0.674395	USD
3	0.027725		United States		0.563360	USD
4	0.584926		United States		0.700511	USD

	Customer Id	Invoice Date	Product Id	Variant Id	Quantity
0	0.440769	2025-03-19 17:27:00	0.063524	0.069303	0.0
1	0.767231	2025-03-24 15:42:00	0.063524	0.069303	0.0
2	0.841846	2025-03-22 18:32:00	0.063524	0.069303	0.0
3	0.496462	2025-03-18 10:51:00	0.063524	0.069303	0.0
4	0.288308	2025-03-22 09:55:00	0.063524	0.069303	0.0

	Subtotal Price	Country_Product_United States - Clogs
0	0.016327	False
1	0.023673	False
2	0.026395	False
3	0.008163	False
4	0.016327	False

	Country_Product_United States - Cycling Shoes
0	False
1	False
2	False
3	False
4	False

	Country_Product_United States - Flip-Flops
0	False
1	False

2	False
3	False
4	False

Country_Product_United States - Gift Card \	
0	False
1	False
2	False
3	False
4	False

Country_Product_United States - Jackets \	
0	False
1	False
2	False
3	False
4	False

Country_Product_United States - Running Shoes \	
0	False
1	False
2	False
3	False
4	False

Country_Product_United States - Sandals \	
0	False
1	False
2	False
3	False
4	False

Country_Product_United States - Tennis Shoes \	
0	False
1	False
2	False
3	False
4	False

Country_Product_United States - Walking Shoes \	
0	False
1	False
2	False
3	False
4	False

Country_Product_United States - Water Shoes	
0	False
1	False
2	False

```
3
4                                     False
                                     False
```

```
[5 rows x 27211 columns]
```

```
df.head()
```

	Order Number	Billing Address	Country	Billing Address	Zip	Currency
0	0.226245		United States		0.884583	USD
1	0.938358		United States		0.504696	USD
2	0.672140		United States		0.674395	USD
3	0.027725		United States		0.563360	USD
4	0.584926		United States		0.700511	USD

	Customer Id	Invoice Date	Product Id	Variant Id	Quantity
0	0.440769	2025-03-19 17:27:00	0.063524	0.069303	0.0
1	0.767231	2025-03-24 15:42:00	0.063524	0.069303	0.0
2	0.841846	2025-03-22 18:32:00	0.063524	0.069303	0.0
3	0.496462	2025-03-18 10:51:00	0.063524	0.069303	0.0
4	0.288308	2025-03-22 09:55:00	0.063524	0.069303	0.0

	Subtotal Price	...	Country_Product_United States	- Clogs
0	0.016327	...		False
1	0.023673	...		False
2	0.026395	...		False
3	0.008163	...		False
4	0.016327	...		False

	Country_Product_United States	- Cycling Shoes
0		False
1		False
2		False
3		False
4		False

	Country_Product_United States	- Flip-Flops
0		False
1		False
2		False

3	False
4	False

	Country_Product_United States - Gift Card \
0	False
1	False
2	False
3	False
4	False

	Country_Product_United States - Jackets \
0	False
1	False
2	False
3	False
4	False

	Country_Product_United States - Running Shoes \
0	False
1	False
2	False
3	False
4	False

	Country_Product_United States - Sandals \
0	False
1	False
2	False
3	False
4	False

	Country_Product_United States - Tennis Shoes \
0	False
1	False
2	False
3	False
4	False

	Country_Product_United States - Walking Shoes \
0	False
1	False
2	False
3	False
4	False

	Country_Product_United States - Water Shoes
0	False
1	False
2	False
3	False



4

False

[5 rows x 27211 columns]

## Step - 8 : Correlation Analysis & Feature Selection


### Correlation Matrix for Numerical Features

```
import seaborn as sns
import matplotlib.pyplot as plt

# Compute correlation matrix only for main numerical columns to avoid
# memory issues
corr_matrix = df[numerical_cols].corr()

# Plot heatmap
plt.figure(figsize=(10, 8))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', fmt=".2f",
            linewidths=0.5)
plt.title("Correlation Matrix of Main Numerical Features")
plt.tight_layout()
plt.show()

C:\Users\moam\AppData\Local\Temp\ipykernel_16228\1999753851.py:11:
UserWarning: Glyph 128279 (\N{LINK SYMBOL}) missing from current font.
    plt.tight_layout()
c:\ProgramData\anaconda3\Lib\site-packages\IPython\core\
pylabtools.py:152: UserWarning: Glyph 128279 (\N{LINK SYMBOL}) missing
from current font.
    fig.canvas.print_figure(bytes_io, **kw)
```



## Identify Highly Correlated Features

```
# Get pairs with correlation > 0.8 (excluding perfect correlation with itself)
threshold = 0.8
high_corr_pairs = []

for i in range(len(corr_matrix.columns)):
    for j in range(i):
        if abs(corr_matrix.iloc[i, j]) > threshold:
            col1 = corr_matrix.columns[i]
            col2 = corr_matrix.columns[j]
            corr_value = corr_matrix.iloc[i, j]
            high_corr_pairs.append((col1, col2, corr_value))

# Show results
for col1, col2, val in high_corr_pairs:
    print(f"□ {col1} and {col2} are highly correlated ({val:.2f})")
```

```
□ Variant Id and Product Id are highly correlated (0.98)
□ Total Price Usd and Subtotal Price are highly correlated (1.00)
□ Total Tax and Subtotal Price are highly correlated (1.00)
□ Total Tax and Total Price Usd are highly correlated (1.00)
```

## Export Cleaned Dataset to CSV

```
# Save as CSV
df.to_csv('Cleaned_Shopify_Sales.csv', index=False)

print("□ Dataset successfully exported!")
```